
02450 Project 1

Report

by Karol Dzitkowski
Marco Becattini

Technical University of Denmark
Department of Applied Mathematics and Computer Science
Introduction to Machine Learning and Data Mining
Tue Herlau
27th September 2014

Abstract

The objective of this report is to apply the methods we learned in the first section of the course on "Data: Feature extraction, and visualization" on letter recognition data set to get a basic understanding of the data prior to the further analysis. This report includes a description of the dataset, detailed explanation of the attributes of the data, data visualizations and a conclusion what we found out about the data set.

Data set

Description

Title Letter Image Recognition Data

Characteristics Multivariate

Number of Instances 20000

Attribute Characteristics Integer

Number of Attributes 17

Missing Values? No

The objective is to identify each of a large number of black-and-white rectangular pixel displays as one of the 26 capital letters in the English alphabet. The character images were based on 20 different fonts and each letter within these 20 fonts was randomly distorted to produce a file of 20,000 unique stimuli. Each stimulus was converted into 16 primitive numerical attributes (statistical moments and edge counts) which were then scaled to fit into a range of integer values from 0 through 15.

Data was obtained from a website:

<http://archive.ics.uci.edu/ml/datasets/Letter+Recognition>

and was a part of an article "Letter Recognition Using Holland-style Adaptive Classifier" available on website:

<http://link.springer.com/article/10.1007/BF00114162>

The research for this article investigated the ability of several variations of Holland-style adaptive classifier systems to learn to correctly guess the letter categories associated with vectors of 16 integer attributes extracted from raster scan images of the letters. The best accuracy obtained was a little over 80%. We typically train on the first 16000 items and then use the resulting model to predict the letter category for the remaining 4000. See the article cited above for more details.

Analysis

- The aim of our future machine learning algorithm working on data set is classification, that is to assign a set of values (our data record) to a capital letter in the English alphabet. It will be probably implemented as some kind of neural network. Classification is a hard problem for a letter recognition, because there is no direct algorithm of doing that, what is the reason for implementing artificial intelligence based on neural network.
- Clustering groups similar characters together, what can be implemented on our data using some sort of N dimensional metric, for example Euclidean or another appropriate similarity measure for records describing letters.
- Association rule discovery will be considered here as a way of predicting a probable value of other attribute values basing on some which are known. We will try to find a set of rules which can tell us a range the other values will be in or the most likely value. For example:

$$\{x = 5, y = 6\} \Rightarrow z \in [-1, 1]$$

or

$$\{x = 1, y = 4, z = 5\} \Rightarrow letter = "a" or "e"$$

- Anomaly detection is used to identify records corresponding to a particular letter with a significant deviation in some or all attributes. It will be helpful to identify unusual patterns of some letters.
- We found no logical reason for applying regression on our data set. However, it could be applied to estimate the expected value of random variable for example for some attribute describing a letter, to predict another probable values of this attribute.

Attributes

In this chapter we will introduce the attributes of the data set and provide an analysis of the data by means of statistical tools.

Description

Each record in the data set contains a letter name and 16 attributes describing a letter from English alphabet. These attributes were scaled linearly to a range of integer values from 0 to 15.

No	Name	Description	Type
1.	lettr	capital letter	discrete/nominal (A to Z)
2.	x-box	horizontal position of box	discrete/ratio (integer)
3.	y-box	vertical position of box	discrete/ratio (integer)
4.	width	width of box	discrete/ratio (integer)
5.	high	height of box	discrete/ratio (integer)
6.	onpix	total # on pixels	discrete/ratio (integer)
7.	x-bar	mean x of on pixels in box	discrete/ratio (integer)
8.	y-bar	mean y of on pixels in box	discrete/ratio (integer)
9.	x2bar	mean x variance	discrete/ratio (integer)
10.	y2bar	mean y variance	discrete/ratio (integer)
11.	xybar	mean x y correlation	discrete/ratio (integer)
12.	x2ybr	mean of $x * x * y$	discrete/ratio (integer)
13.	xy2br	mean of $x * y * y$	discrete/ratio (integer)
14.	x-ege	mean edge count left to right	discrete/ratio (integer)
15.	xegvy	correlation of x-ege with y	discrete/ratio (integer)
16.	y-ege	mean edge count bottom to top	discrete/ratio (integer)
17.	yegvx	correlation of y-ege with x	discrete/ratio (integer)

Table 1.1: Attribute Information

Analysis

Data set is describing 26 classes - letters in English alphabet. The distribution of all 20000 records to the classes is nearly uniform:

789 A	766 B	736 C	805 D	768 E	775 F	773 G
734 H	755 I	747 J	739 K	761 L	792 M	783 N
753 O	803 P	783 Q	758 R	748 S	796 T	813 U
764 V	752 W	787 X	786 Y	734 Z		

Table 1.2: Class Distribution - Letters

For each attribute, we computed basic summary statistics: mean, variance, median, standard deviation, as given in the table below:

	Mean	Variance	Median	Std
x-box	4.0236	3.6604	4.0000	1.9132
y-box	7.0355	10.9201	7.0000	3.3046
width	5.1219	4.0585	5.0000	2.0146
high	5.3724	5.1139	6.0000	2.2614
onpix	3.5059	4.7981	3.0000	2.1905
x-bar	6.8976	4.1048	7.0000	2.0260
y-bar	7.5004	5.4073	7.0000	2.3254
x2bar	4.6286	7.2898	4.0000	2.7000
y2bar	5.1787	5.6683	5.0000	2.3808
xybar	8.2820	6.1925	8.0000	2.4885
x2ybr	6.4540	6.9225	6.0000	2.6311
xy2br	7.9290	4.3290	8.0000	2.0806
x-ege	3.0461	5.4407	3.0000	2.3325
xegvy	8.3389	2.3924	8.0000	1.5467
y-ege	3.6917	6.5899	3.0000	2.5671
yegvx	7.8012	2.6162	8.0000	1.6175

Table 1.3: Attribute Information

All attributes have a range from 0 to 15, because they were normalized in that way and there is no missing or damaged values. Each record (set of attributes) is describing one appearance of a letter. In the next section we will consider visualization of that data.

Visualization

In this section we provide an analysis of the data in terms of their graphical representation and visualization.

Histograms

We plotted a histogram for each attribute to investigate its distribution. Each histogram shows how many records had a value of attribute in a range of particular bucket.

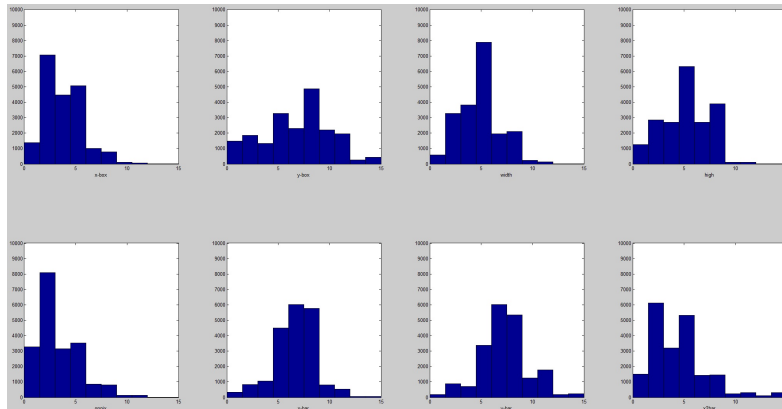


Figure 1.1: Histograms for first attributes 0-7

We can notice that most of the attributes are normally distributed.




Figure 1.2: Histograms for first attributes 8-15

Box plots

We created box plot of our data in order to show the distribution of the attribute values. We did it for all the records in dataset:




Figure 1.3: Box plot for all records of the data set

Also for first 3 classes (letter) to show how attributes differ from one class to another:




Figure 1.4: Box plot for classes A,B,C of the data set

Principal Component Analysis

We applied principal component analysis (PCA) on our data set to convert a set of our set of possibly correlated attributes into a set of values of linearly uncorrelated variables called principal components. We normalized the records in our data set by subtracting the mean and dividing by standard deviation for each attribute. After that, we applied single value decomposition to calculate the variation for every principal component (PC).

Variance

Using a PCA technique we noticed that about 28% of the variation is caused by first PC, about 15% by a second one, and about 12% by the 3rd. It means that first 3 principal components generate more than 50% of the whole variance.

The cumulative plot shows that over 9 principal components are needed to explain more than 90% of the variation.




Figure 1.5: Variance represented by each principal component




Figure 1.6: Cumulative variance for principal components

PCA - visualization

Because of the fact that first 3 PC generate more than 50% of the variance, we decided to visualize records of our data set as a 3D points constructed of first 3 PC values. Picture above shows a PCA visualization for first 3 PC of




Figure 1.7: Classes A,B and C represented by first 3 PC

whole data set. In the picture below however, we can see a visualization only for first three letters. We can see that it is possible with some probability to distinguish letters even using only first 3 PC's.




Figure 1.8: All the data set represented by first 3 PC