
02450 Project 2

Report

by Karol Dzitkowski
Marco Becattini

Technical University of Denmark
Department of Applied Mathematics and Computer Science
Introduction to Machine Learning and Data Mining
Tue Herlau
1st November 2014

Abstract

Regression

Classification

0.1 Problem

The classification problem for our data we have chosen to solve is predicting a letter from the english alphabet (26 classes) using its attributes. Each record (letter) has 16 integer attributes in a range of 0 to 15.

0.2 Methods and parameters

We solved this problem using KNN, ANN and Naive Bayes methods. For every method we implemented two levels of cross-validation for estimating optimal parameters. For KNN it was a value of K , in ANN a number of nodes in the hidden layer, and in Naive Bayes method we estimated a number of features (best) used.

K-Nearest Neighbors

KNN algorithm predictions are very accurate for various values of K , but it appears that lower values are better. The lowest error rate is for $k = 1$ despite of having no duplicates in our dataset. See figure 0.1.

Naural Network

For an artificial neural network method we use a perceptron with one hidden layer which uses softmax function on the output layer. We are optimizing number of nodes in a hidden layer. The lowest error rate occurred with a hidden nodes count of approximately 1.5 times a number of attributes, so about 26. This optimal number varies as the training subset changes in cross-validation, but still is close to that number. In the figure 0.2 we can see an average of error rates for various hidden neurons count.

Naive Bayes

For Naive Bayes method we tried to estimate optimal features set (number of features), however no matter of what classifier we used for feature selection




Figure 0.1: K-Nearest Neighbors - K vs error




Figure 0.2: Neural Network - hidden nodes vs error

using recursive feature elimination with cross-validation, the best accuracy was obtained with selecting all features. In a plot (figure 0.3) we can see how accuracy of the method depends on number of features.




Figure 0.3: Naive Bayes - features vs accuracy

0.3 Results

For all methods, the error rate of predicting a letter basing of its attributes is below 50%, what is quite a good result for such a problem. The best results we got using K-Nearest Neighbors algorithm which with euclidesian mertic obtained less than 5% of bad answers. Out error rate is calculated as:

$$\text{error_rate} = (\text{predicted_letters} \neq \text{actual_letters}) / \text{number_of_tests}$$

Below (figure 0.5) we can see a plot describing how good letters are predicted by KNN method. From 5000 points used only a few are away from $x = y$ line. Most of the points are on that line, which means that letters are predicted correctly.

Moreover, fitting our methods for all of the data (20000 records) we obtained on average 4% of an error rate for KNN and 35%, 44% for Neural Network and Naive Bayes accordingly. We used K-fold cross-validation using all that data which forced us to use parallelism computing every fold in a separate thread.




Figure 0.4: Performance of methods in outer cross-validation loop




Figure 0.5: Predicted values vs actual values (KNN)

0.4 Comparison

The boxplot below (figure 0.6) shows that there are great differences between a performance of used methods. We used also a paired t-test to compare

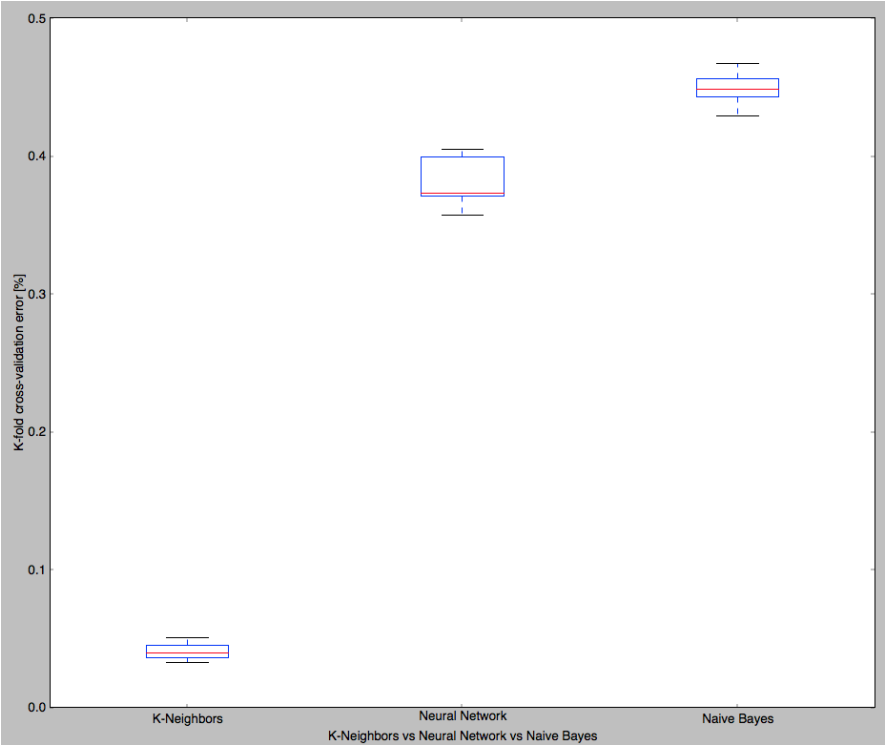


Figure 0.6: Cross-validation errors for all methods

selected methods with each other and additionally with a fake predictor, which classifies all outputs to be the largest class in the training data:

Compared alg.	t-test value	p-value	significantly
ANN vs Naive Bayes	-10.3	0.0	True
ANN vs KNN	57.8	0.0	True
Naive Bayes vs KNN	100.5	0.0	True
Biggest Class vs ANN	-100.4	0.0	True
Biggest Class vs KNN	-380.7	0.0	True
Biggest Class vs Naive Bayes	-133.3	0.0	True

The results shows us that all of the methods are significantly different in case of performance.

0.5 Related work

This dataset was used for letter recognition (classification) using Holland-style Adaptive Classifiers. The article is available on the following link:

http://download.springer.com/static/pdf/733/art%253A10.1007%252F00114162.pdf?auth66=1414859976_6267eff4e2e0779ac8c3bd5c7c57b61c&ext=.pdf The

percentage of correct identifications they achieved varies between 50% to 80% depending on the settings, where the highest score is obtained with the most computationally demanding settings. Referring to that article we can assume that our results are successful.

Conclusions